

A Pipeline for Tailored Sampling for Progressive Visual Analytics

Marius Högräfer, Jakob Burkhardt, Hans-Jörg Schulz

Aarhus University, Denmark

Rome, 13th June 2022



AARHUS
UNIVERSITY

Contributions

1 A technique

- Propose a tailorable sampling pipeline for PVA

2 A demonstrative use case

- Tailor the progressive sampling

3 A tool

- ProSample, which allows comparing two pipelines side-by-side

Primer on Progressive Visual Analytics

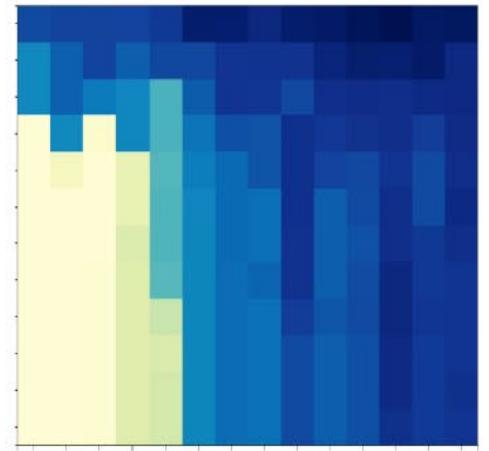
- Analysis on large data takes too long → **Not interactive!**
- Split the data into smaller chunks
- Enable interactive analysis on **early, partial** results

- Analyst gets to see the data they are interested much earlier
- Bring the Human "back into the loop"

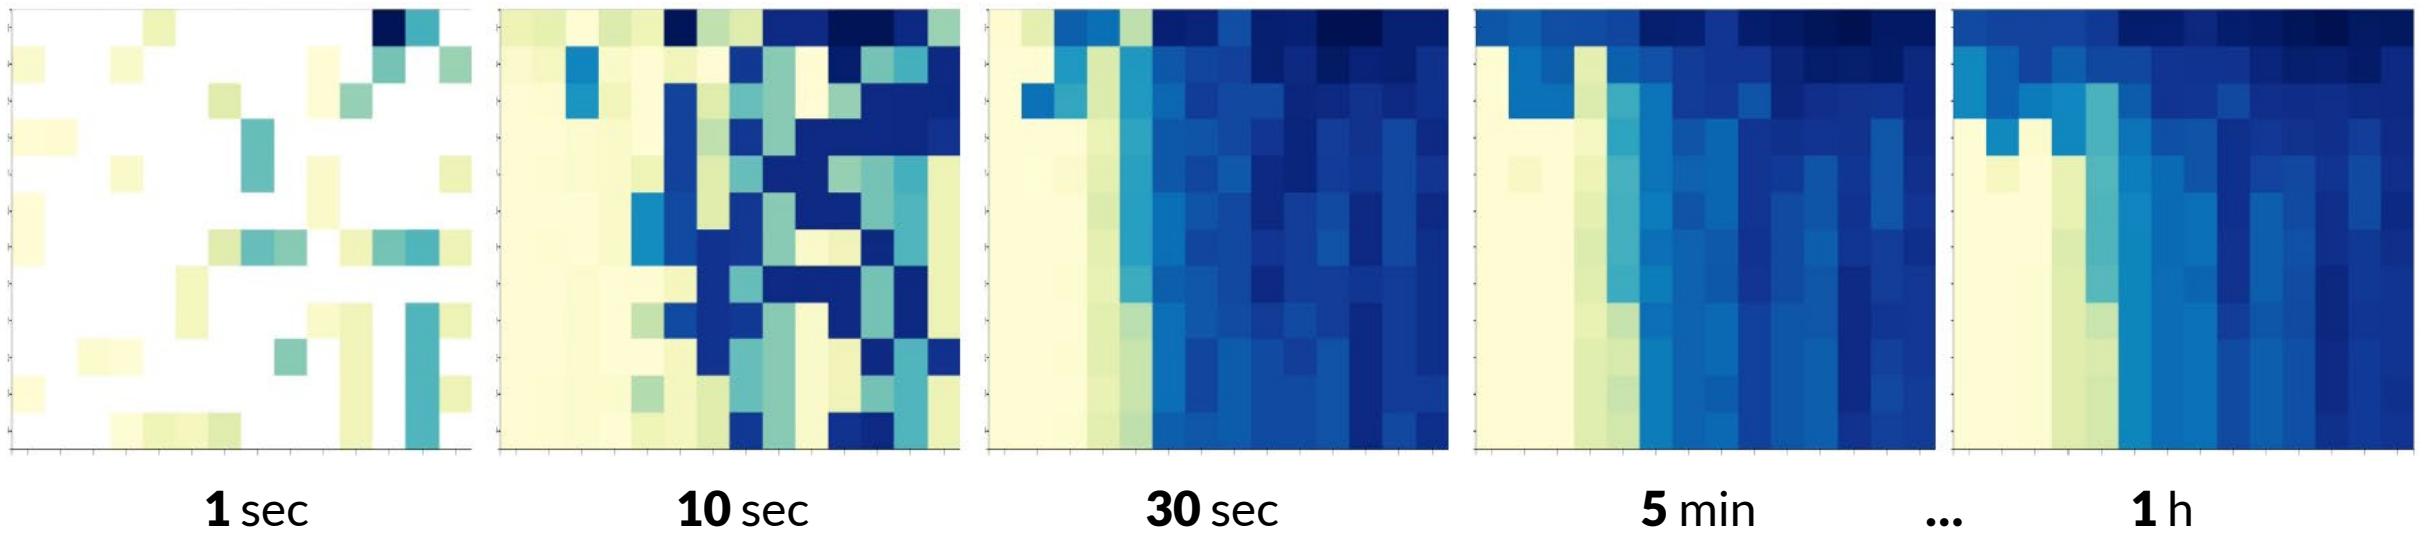
Standard approach:

Launch analysis...

1h later (or more!)

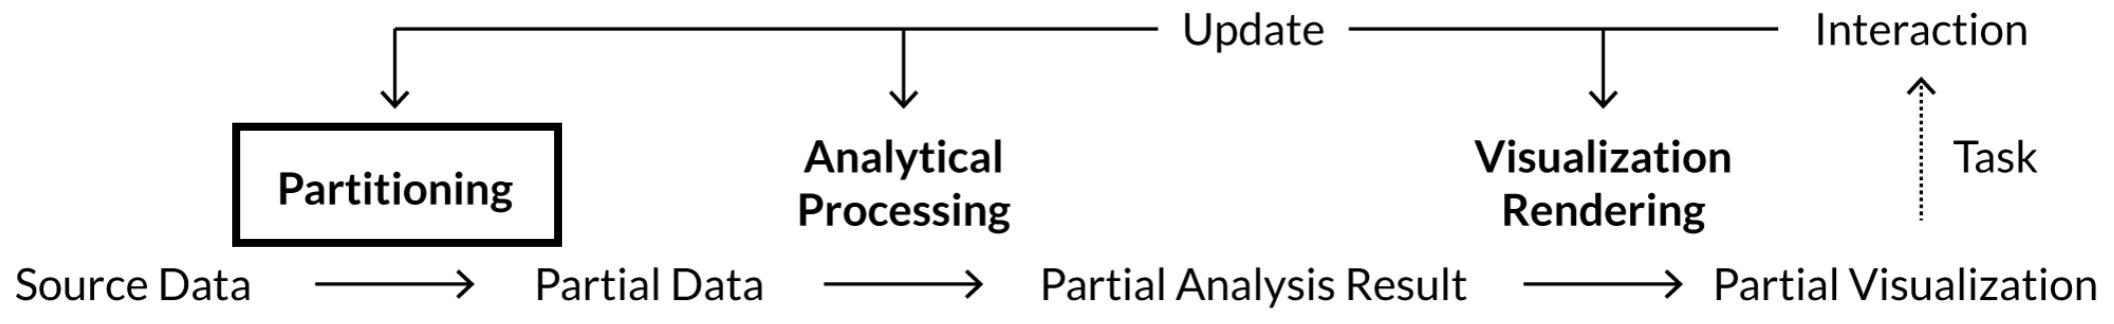


PVA approach:



Sampling in PVA

- Fundamental to PVA: **Data arrives in chunks**
- The first step in the process:



Adapted from [Li+Ma, 2020]

Challenge

In-progress visualization should be **representative** of the final result

- What makes a chunk representative?
- What order should the data arrive in?

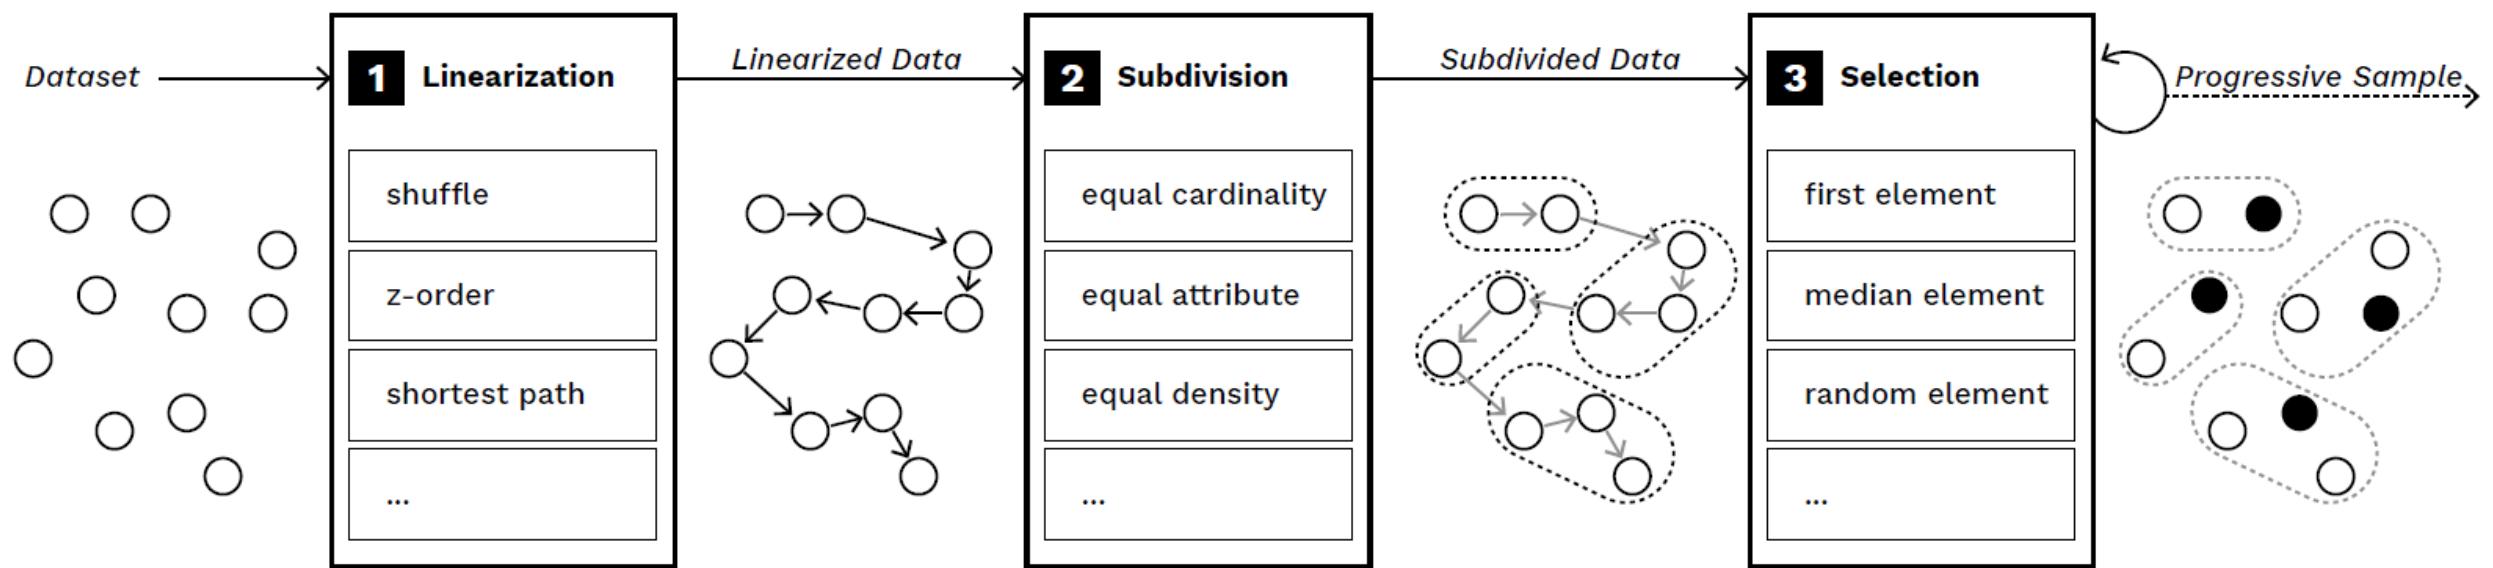
Depends on the analysis scenario!
(see Related Work)

Background

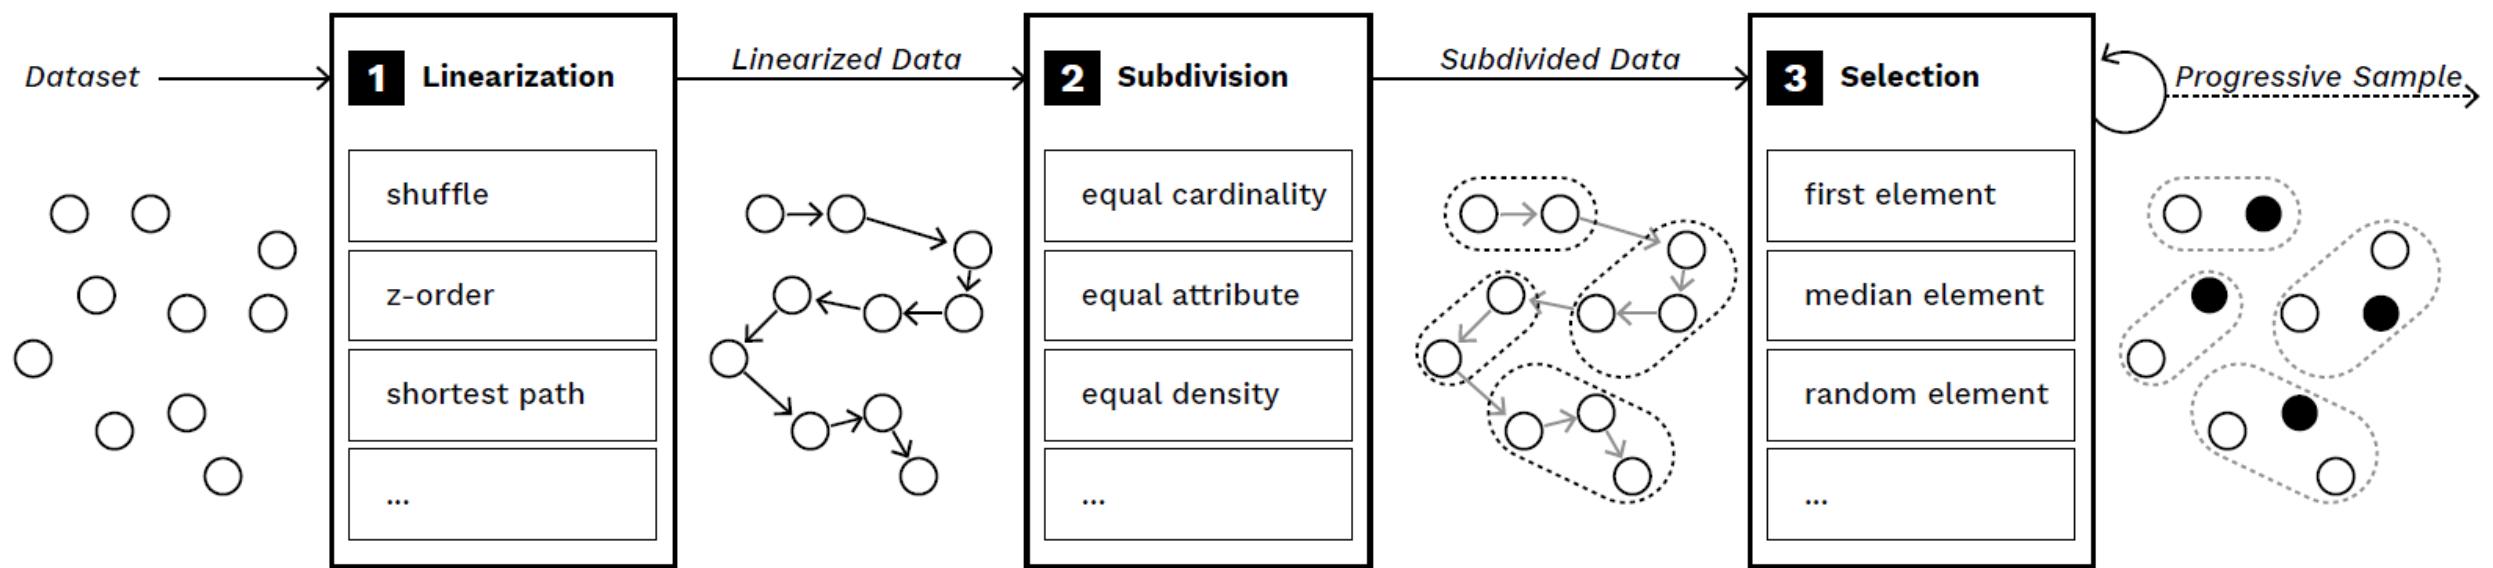
- What to do, when no dedicated algorithm exists?
- Fallback is random sampling ("one-size-fits-all")
 - Can produce visual artifacts on some visualizations [Zheng et al. 2017]
 - Poor fit for tasks like outlier detection [Chen et al. 2022]

Our idea: break up sampling process to modularize it
→ allows to **tailor** it to analysis scenarios

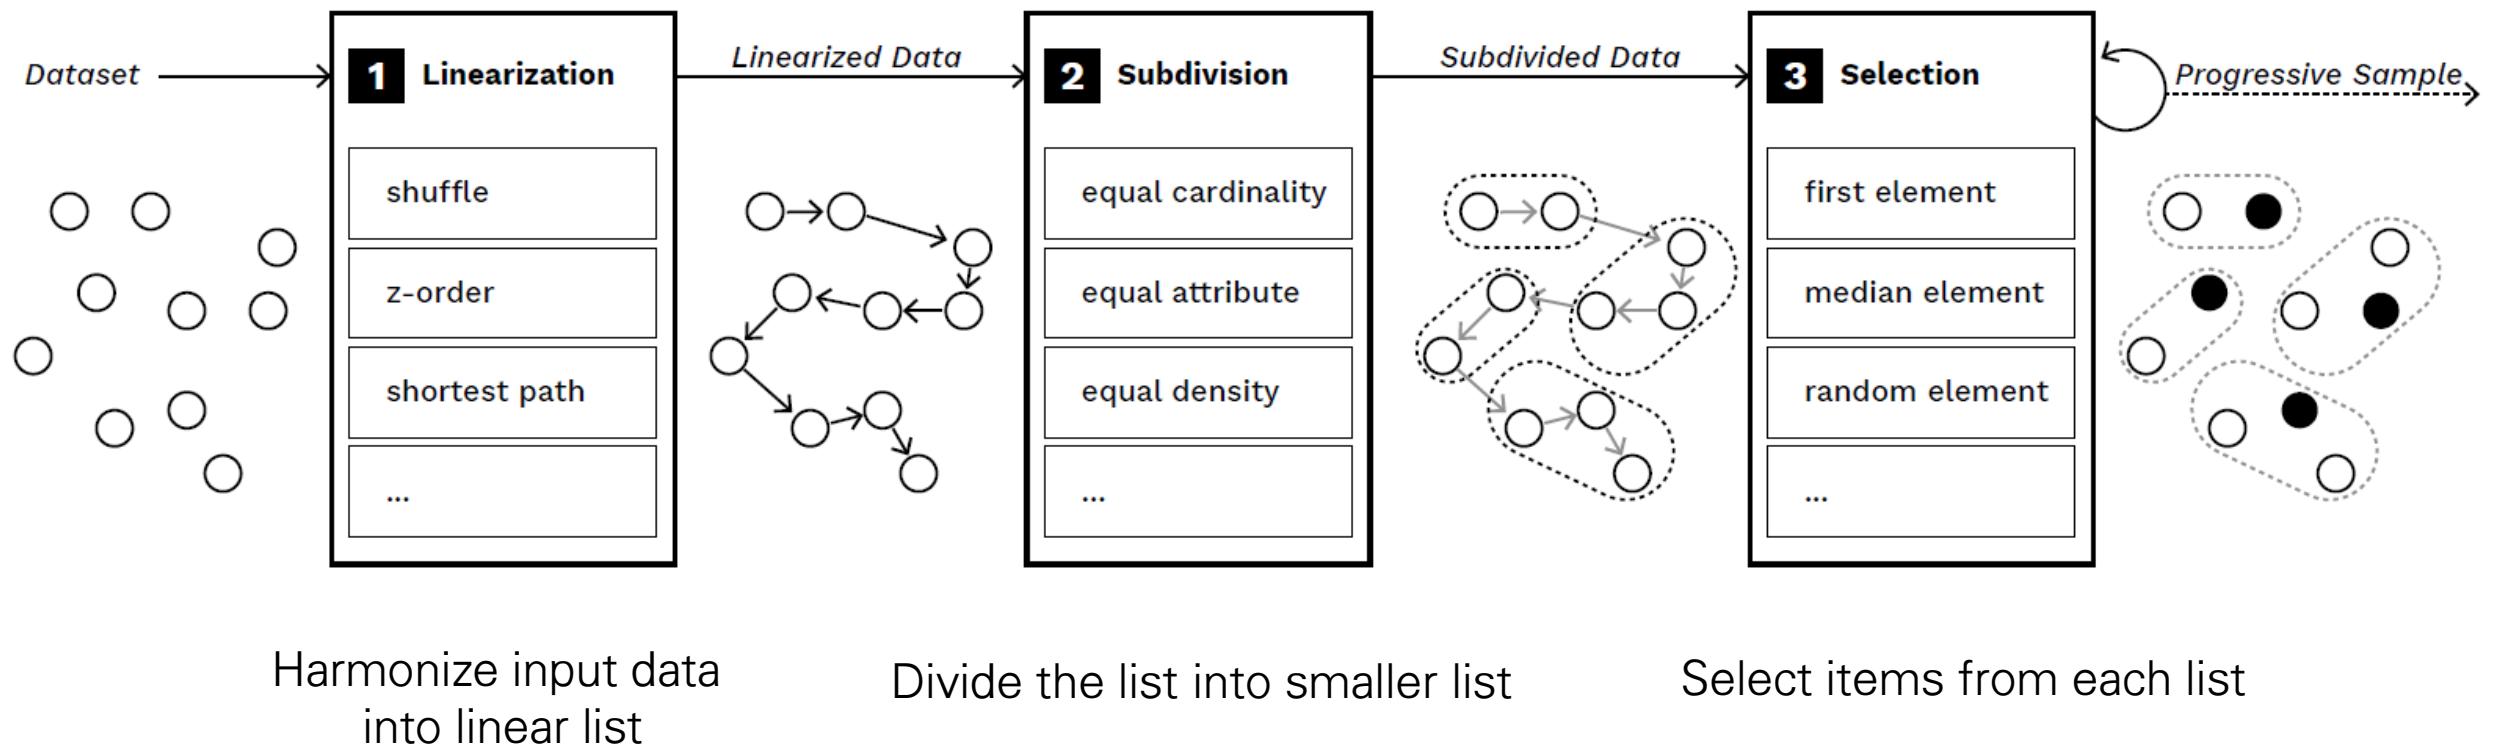
A tailorable Sampling Pipeline for PVA



A tailorable Sampling Pipeline for PVA

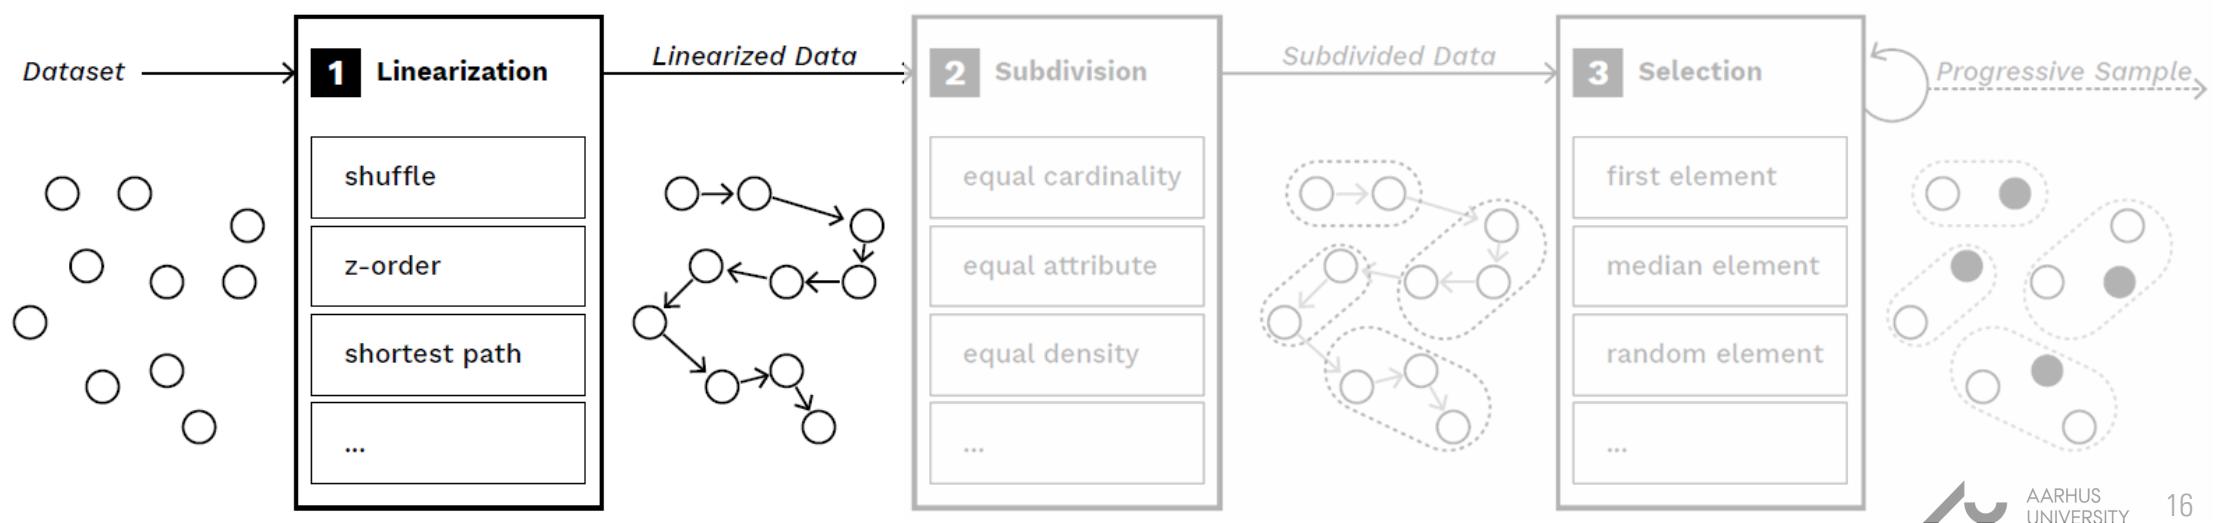


How does the pipeline work?

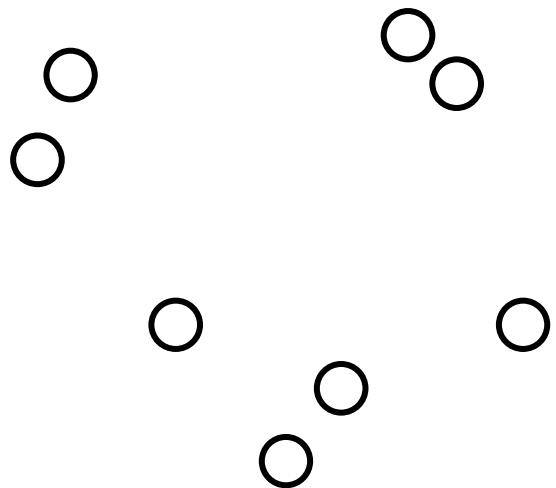


Linearization

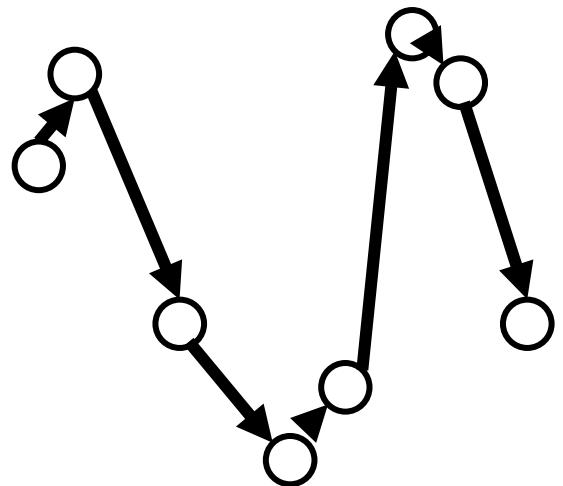
- Input: Dataset
- Output: List
- Tailor to **characteristics of the dataset**
- Based on data structure



Linearization - Example

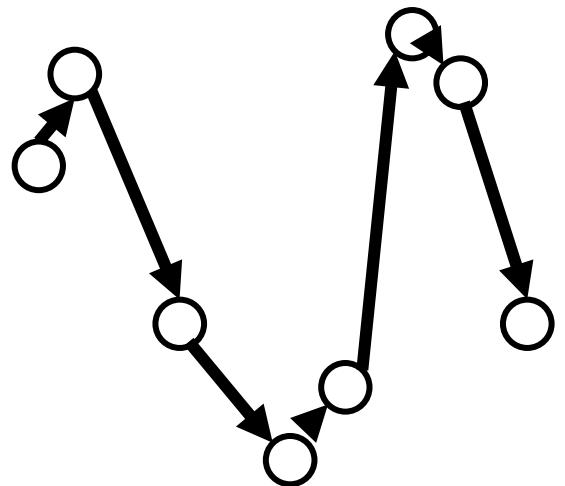


Linearization - Example

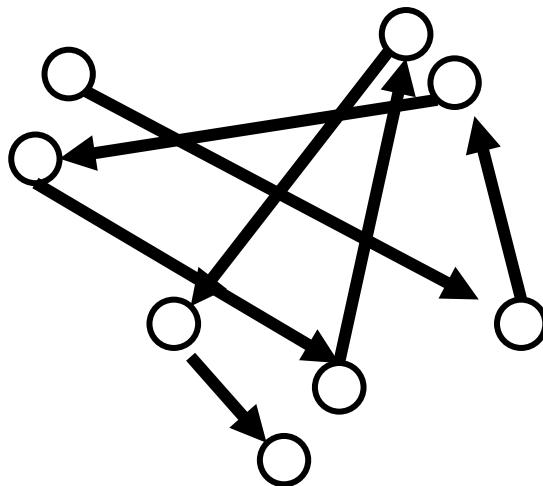


Sort by attribute

Linearization - Example

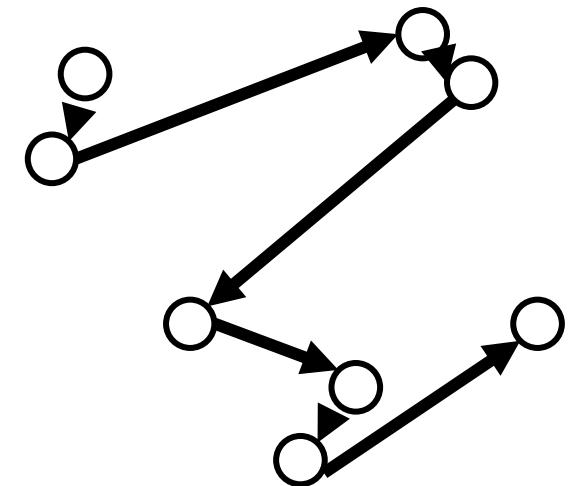


Sort by attribute



Shuffle

[Badam et al. 2013]

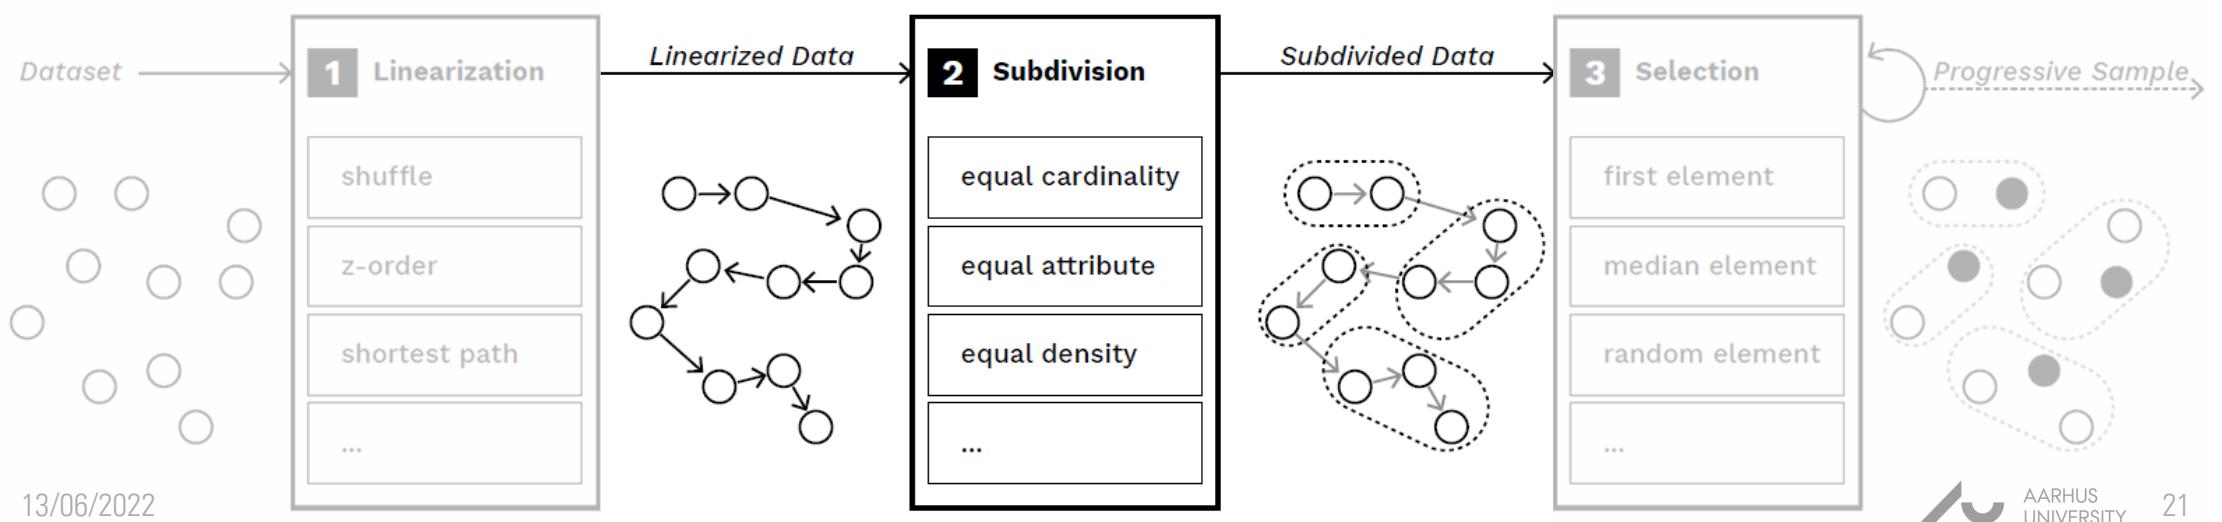


Z-order

[Zhou et al. 2020]

Subdivision

- Input: List of data items
- Output: Groups of lists of data items
- Tailor to the **analysis task**
- Based on data attributes



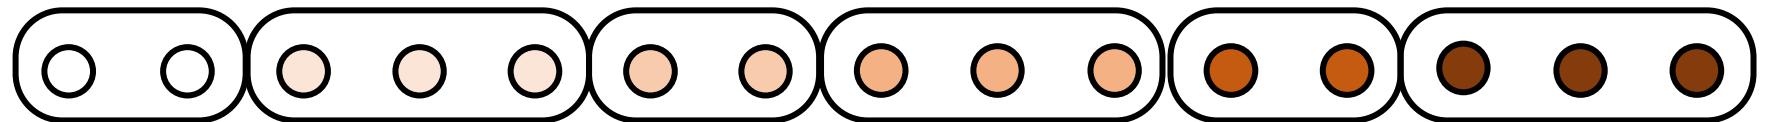
Subdivision - Example

Subdivision - Example

Cardinality:



Attribute:

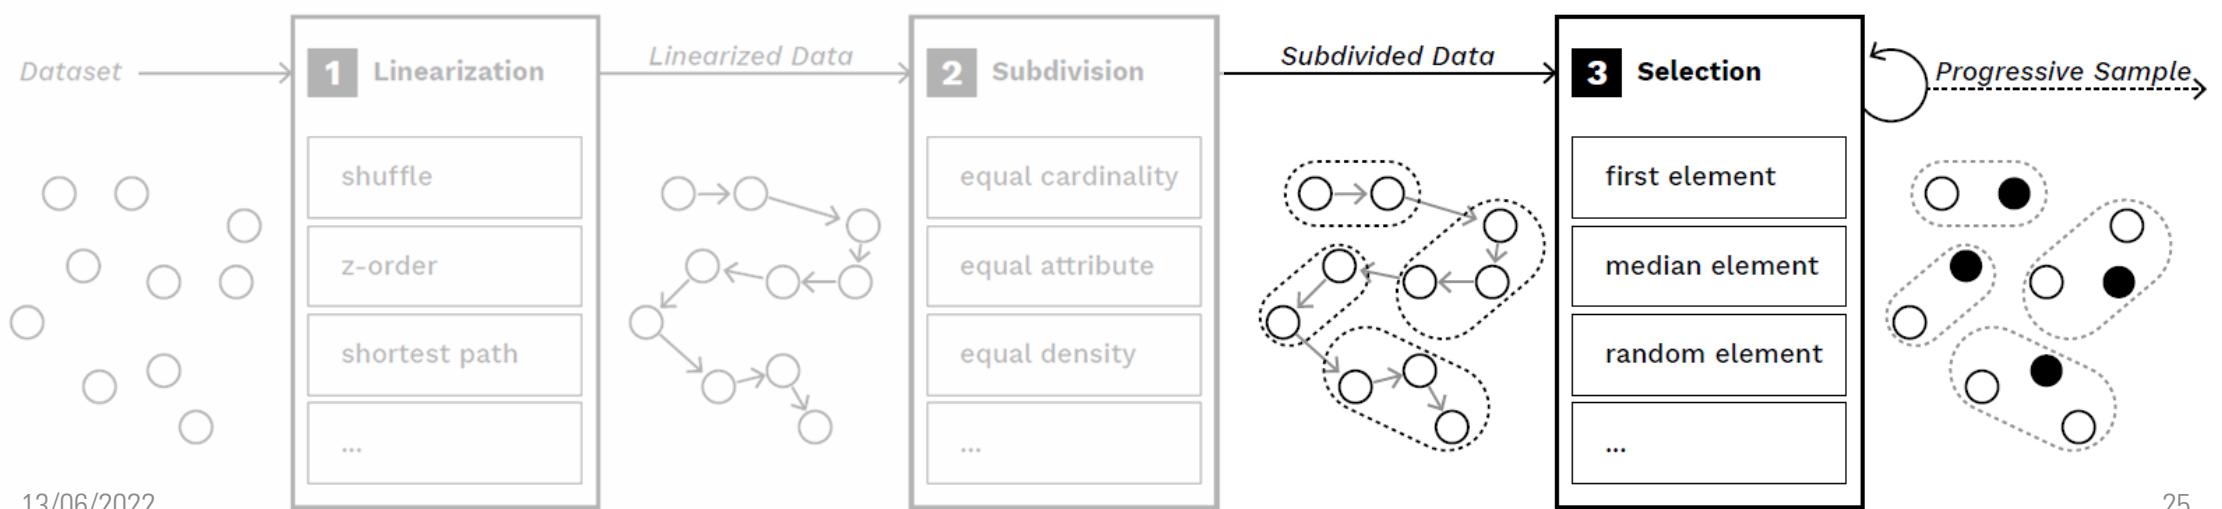


Distance:



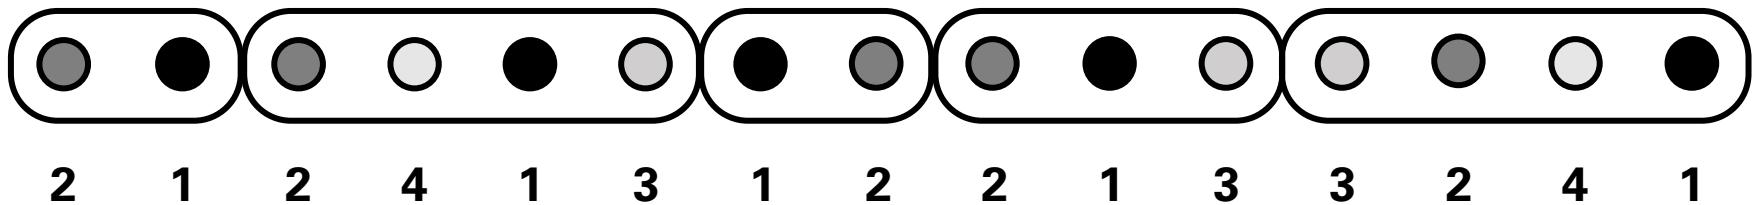
Selection

- Input: Groups of data lists
- Output: Partitions of the dataset
- Tailor to the **user interest** [Micallef et al. 2019]
- Based on the desired order of the data



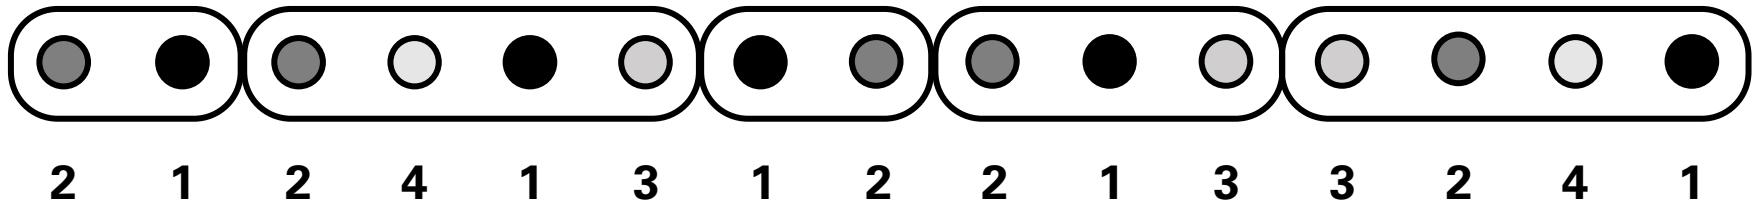
Selection – Example

Random:

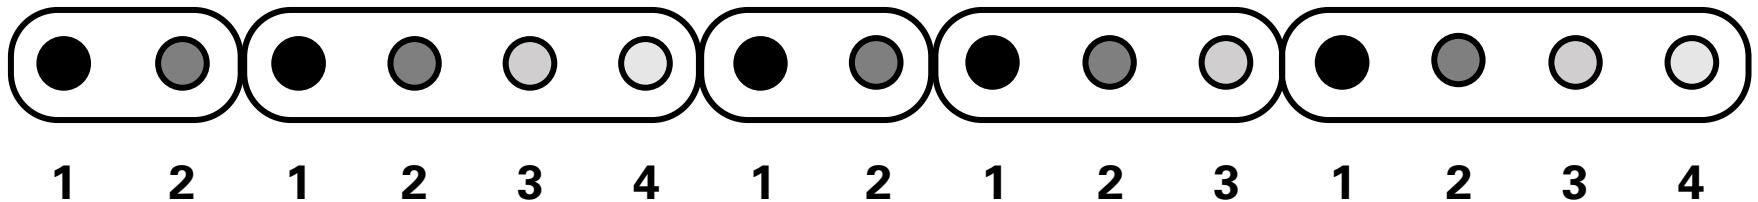


Selection – Example

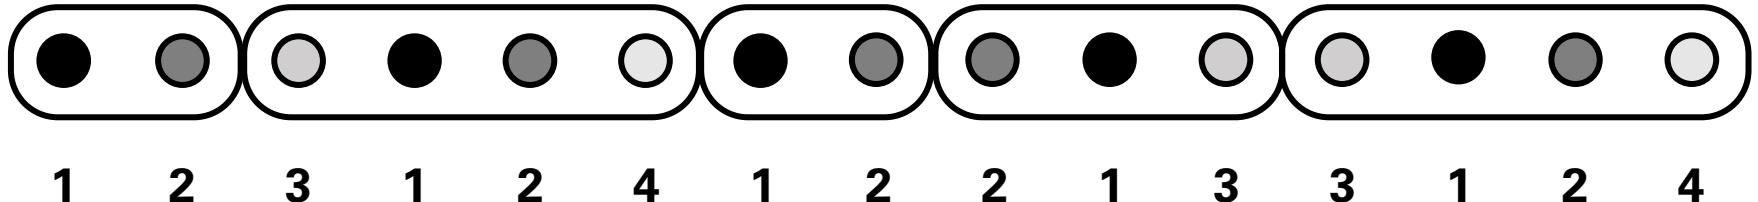
Random:



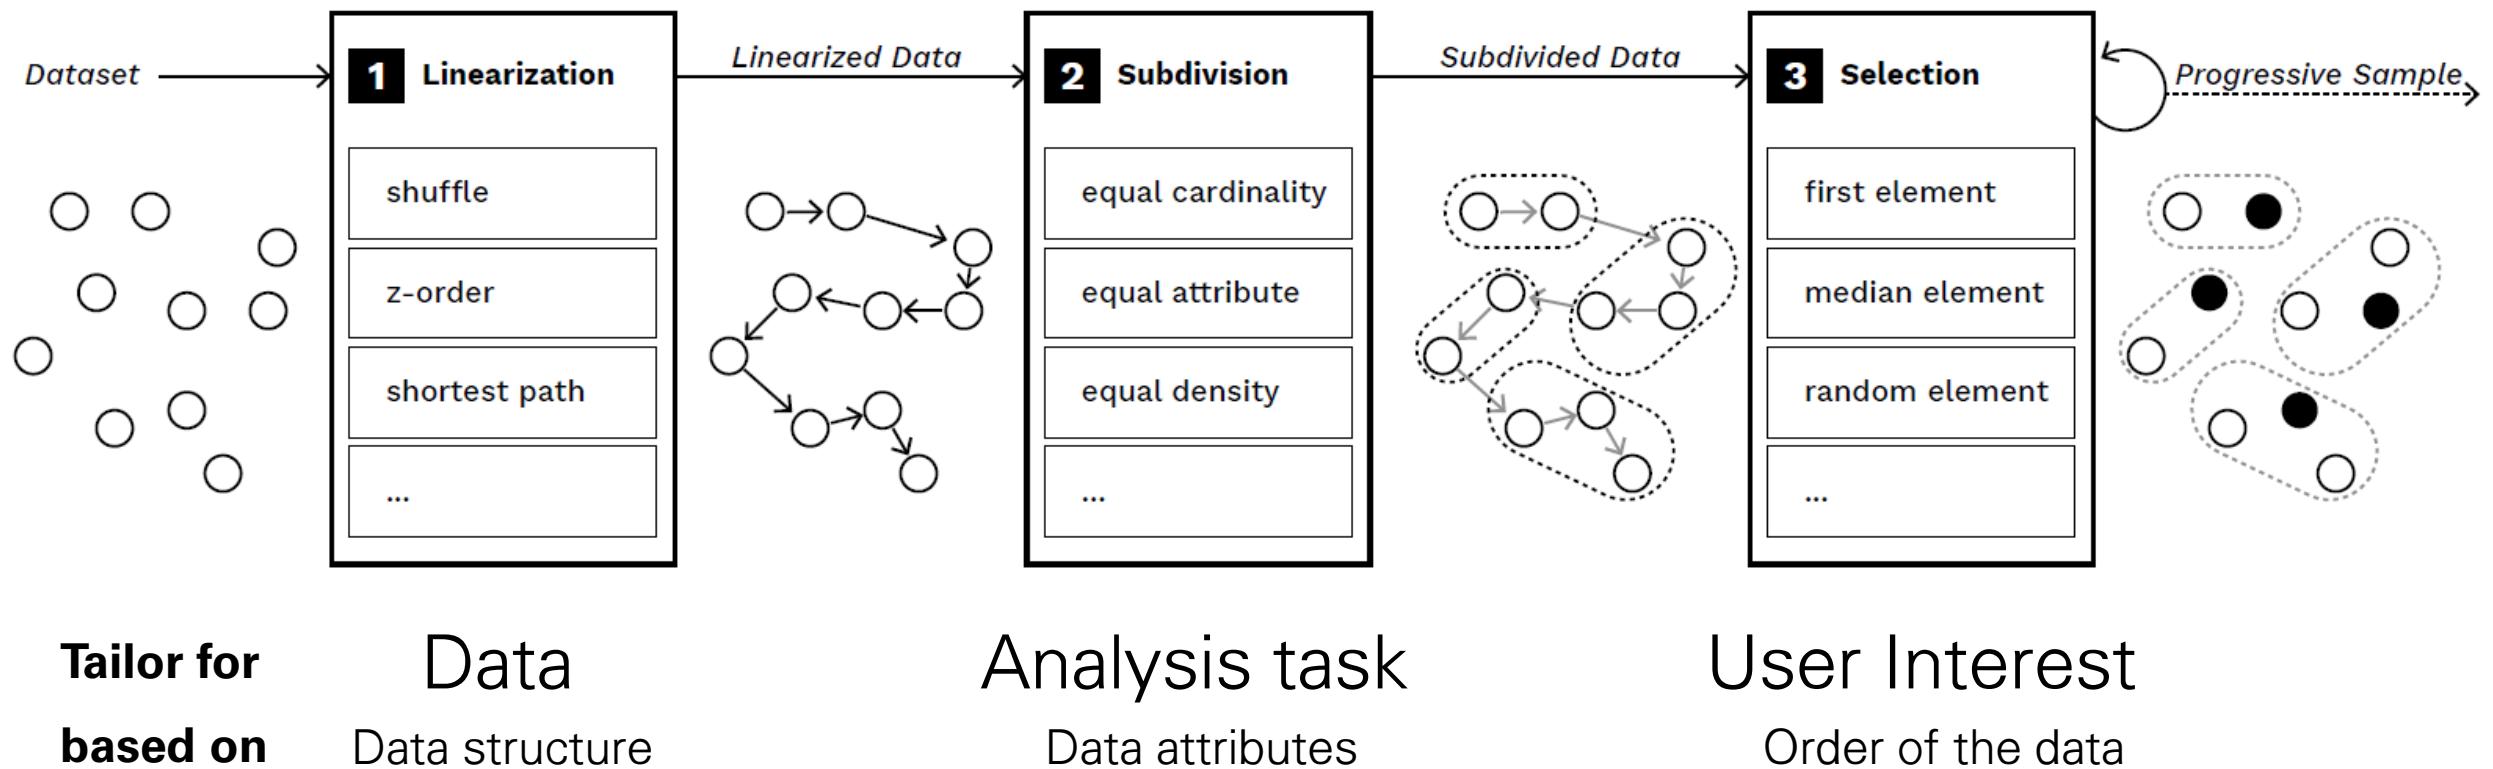
First:



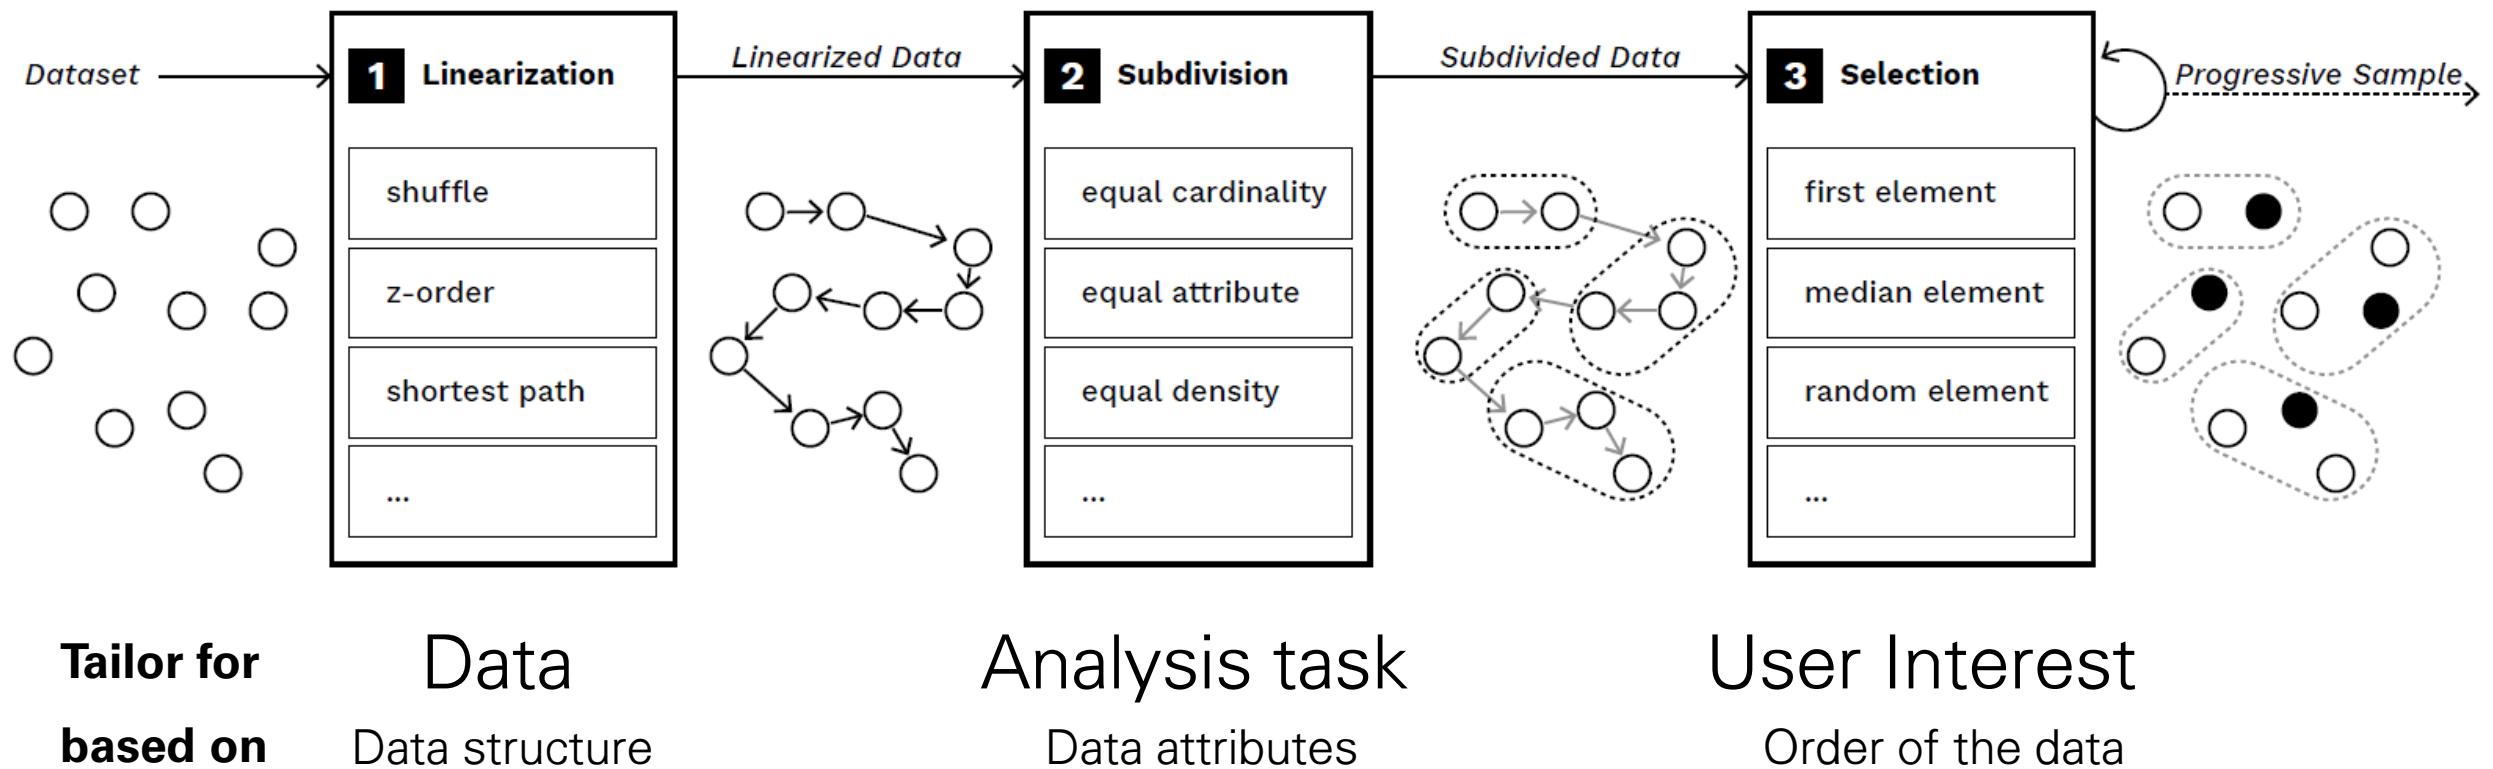
Median:



Tailoring the Sampling



Tailoring the Sampling



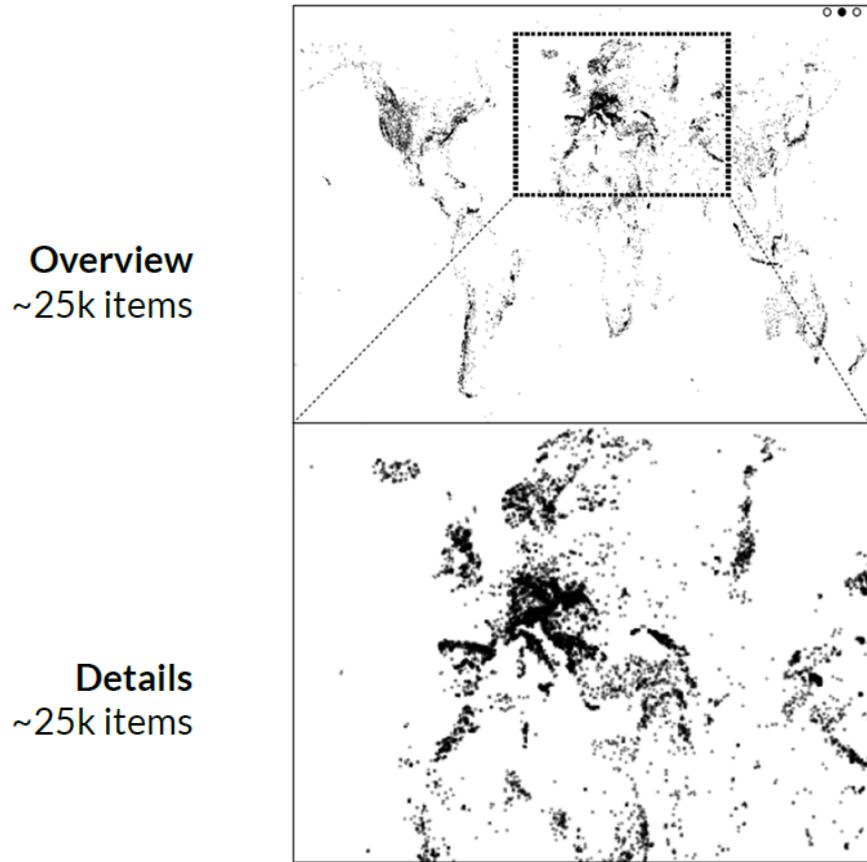
Using the pipeline

Geospatial dataset ("mountain peaks" from OpenStreetMaps)

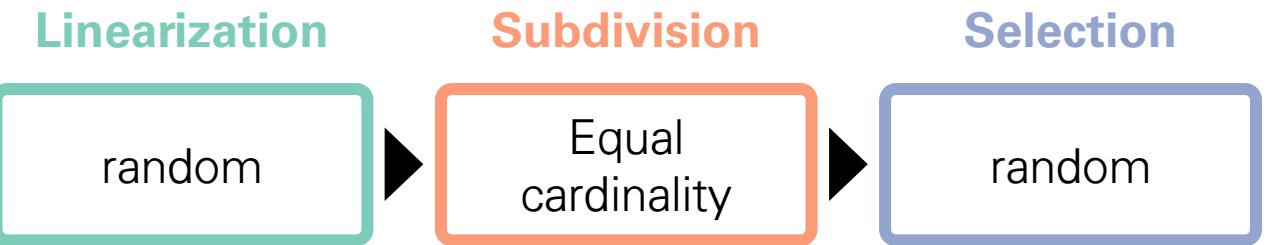
Longitude, latitude, altitude of each peak.

Explore altitude distribution in the dataset.

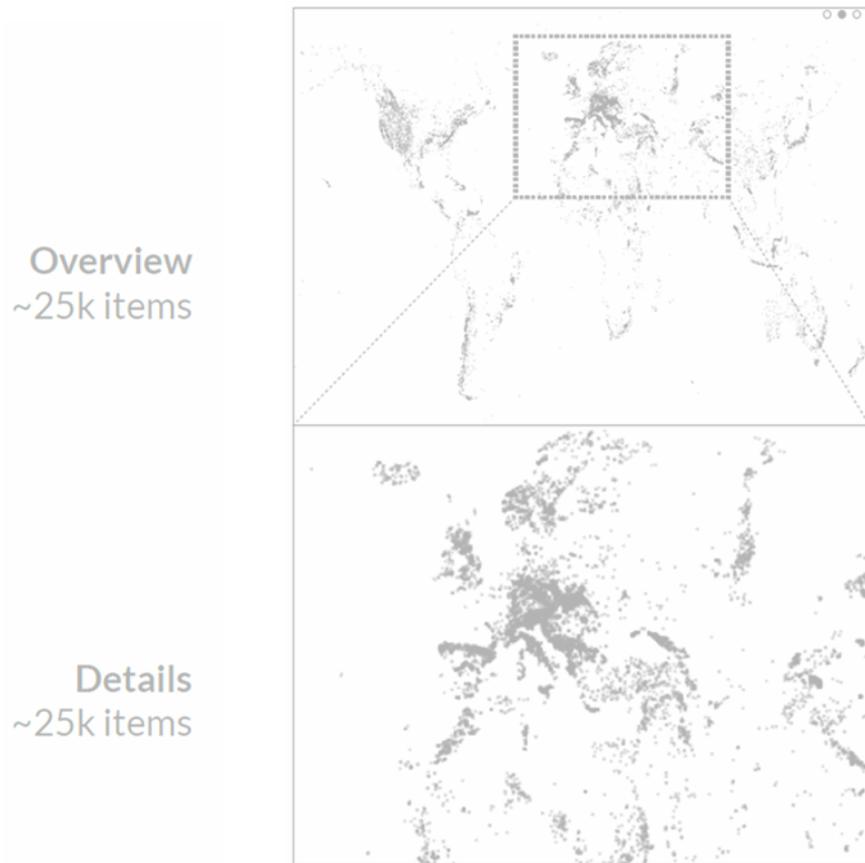
Explore the data



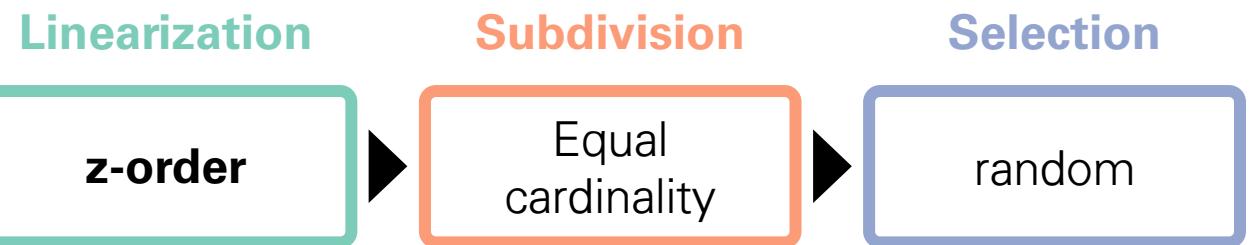
- Develop a "feel" for the dataset
- Make little assumptions about the data in the beginning



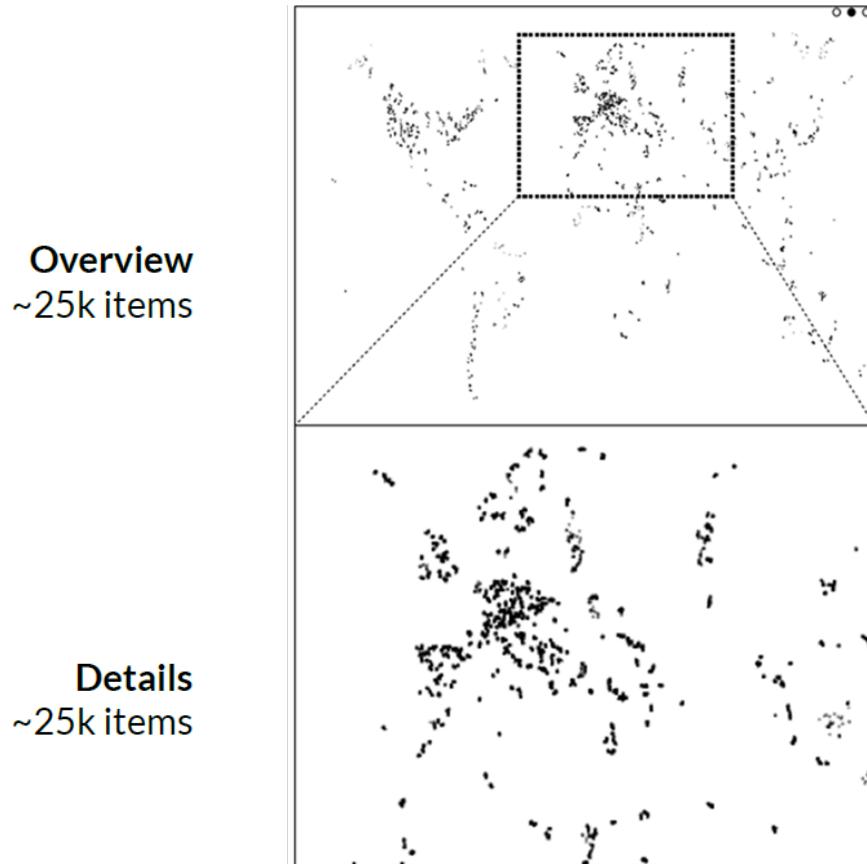
Tailor for Data



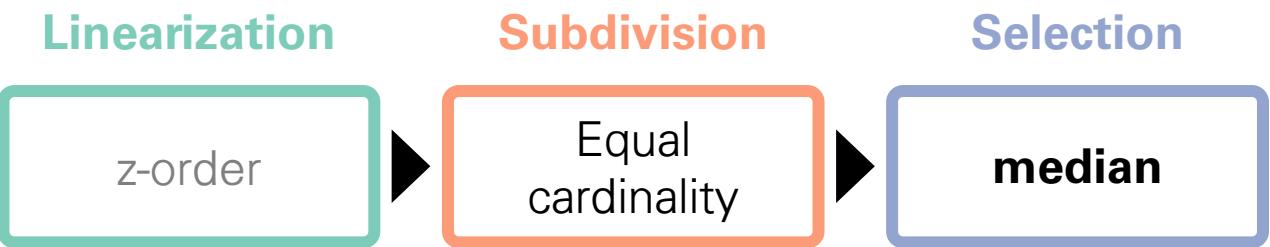
- Develop a "feel" for the dataset
- Make little assumptions about the data in the beginning
- Preserve density+outliers
→ tailor via **Linearization**



Explore average altitude

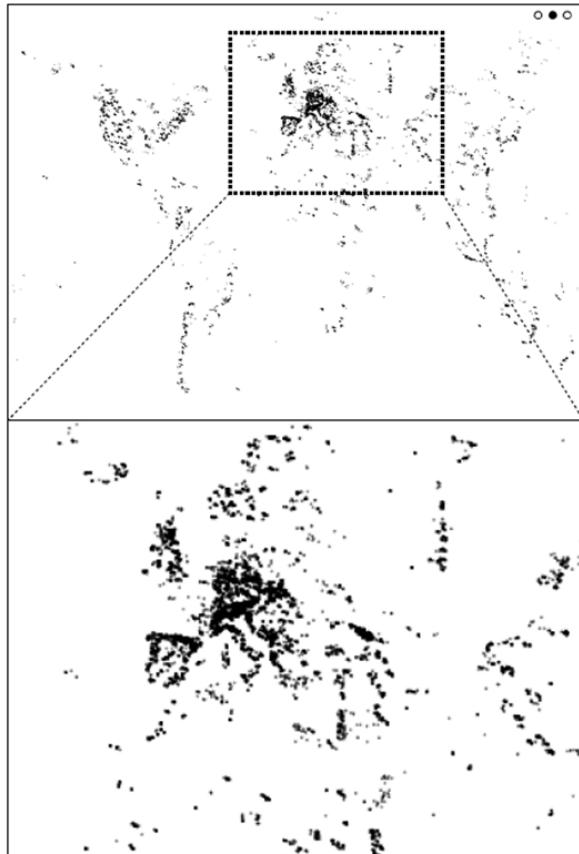


- What is the average altitude in a region?
- Change in user interest
→ tailor via **Selection**



Maintain outliers

Overview
~25k items



- Get distribution of altitude
- Change in user interest
→ Tailor via **Selection**

Linearization

z-order

Subdivision

Equal cardinality

Selection

random

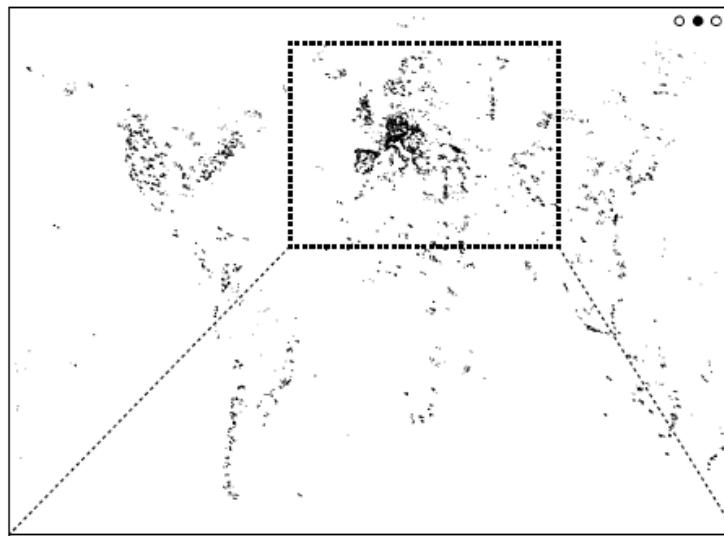
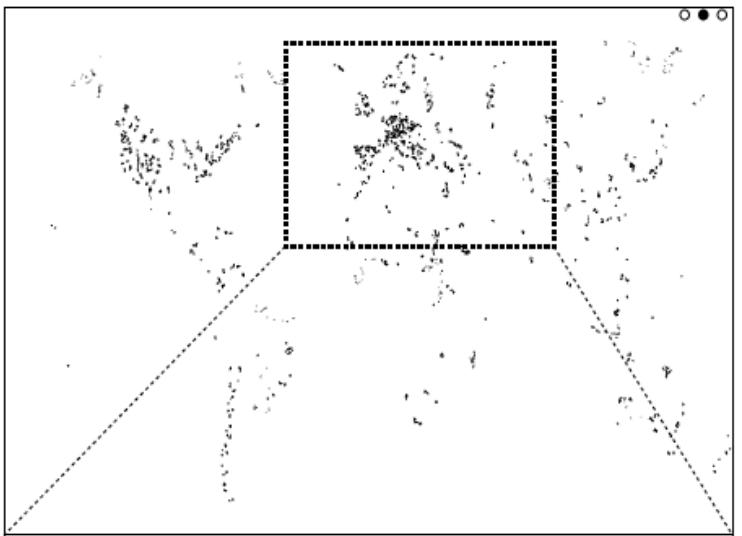
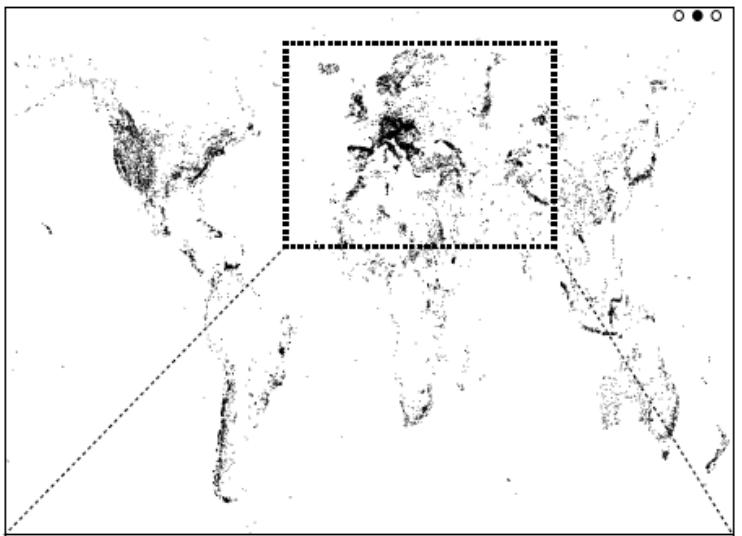
Pipeline

random → equal cardinality → random

z-order → equal cardinality → median

z-order → equal cardinality → random

Overview
~25k items



Details
~25k items



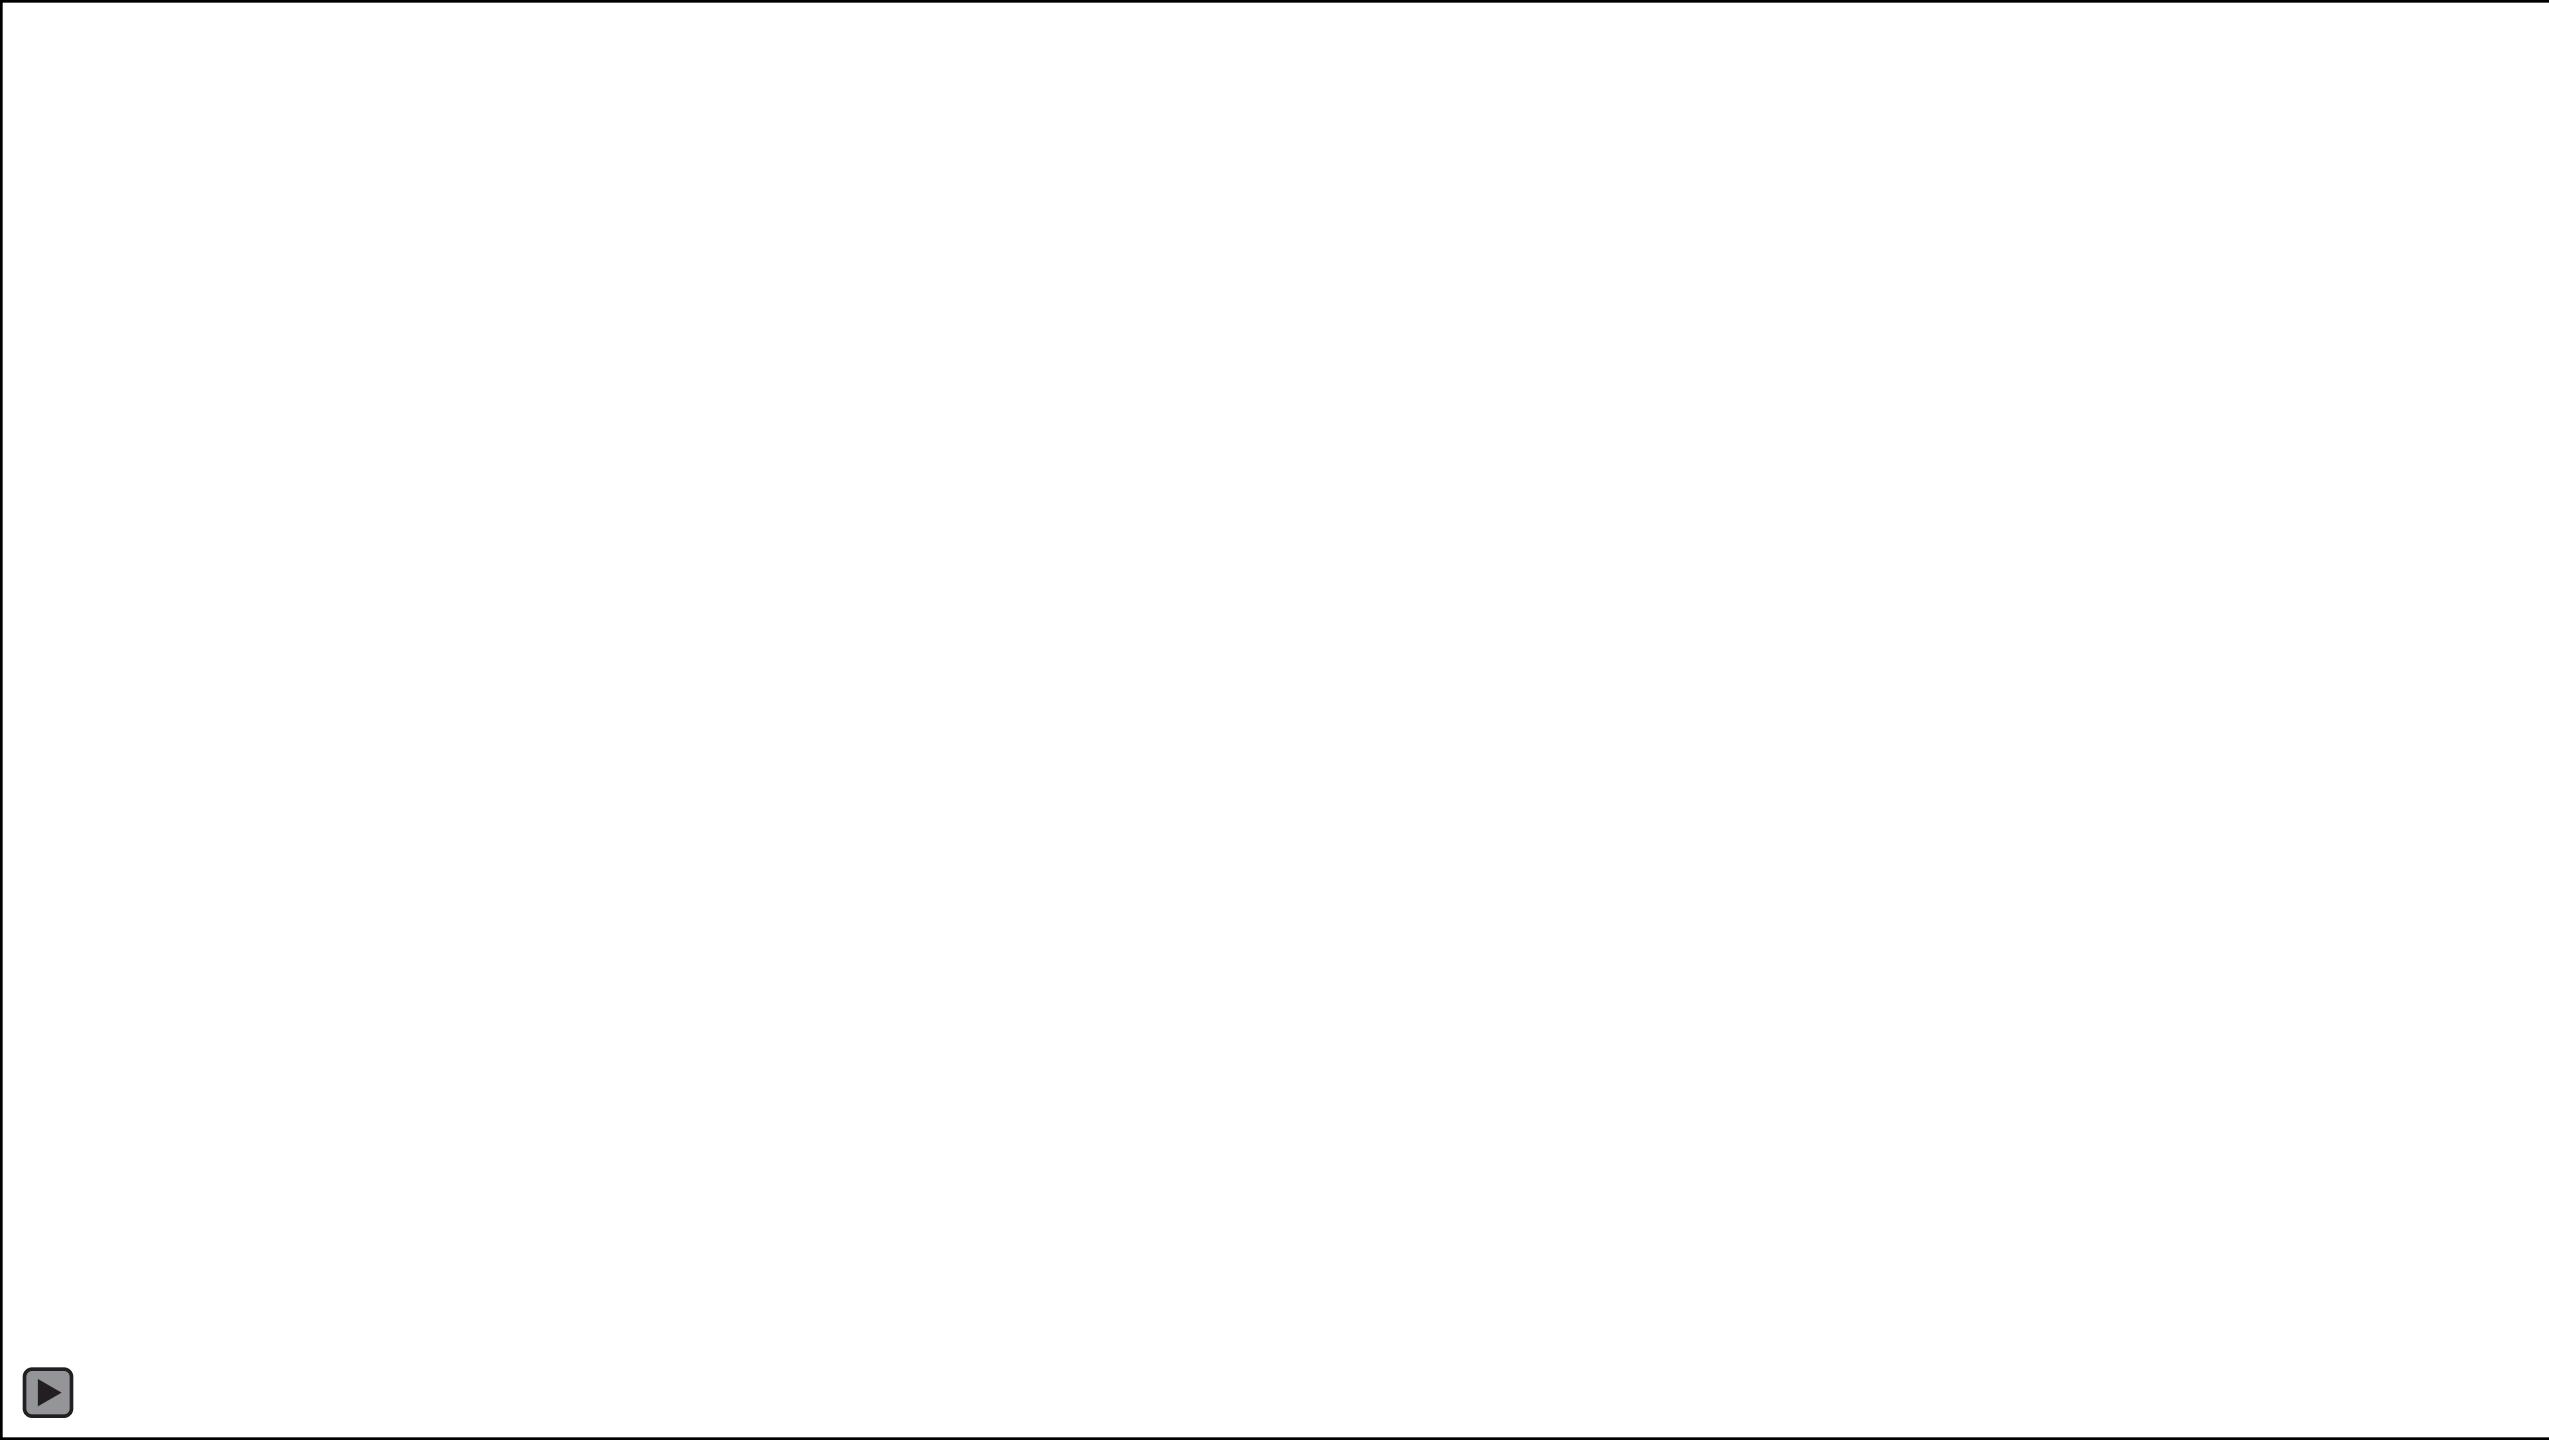
Overview

Median Altitude

Altitude distribution

Comparing Pipelines with ProSample





Future Work

- Develop actionable **guidelines** for pipeline configurations
 - Based on runtime, utility, parameters ...
- Explore **open questions**:
 - Where in the PVA process can we position which sampling pipeline?
 - How to combine with steering approaches? How can we prioritize certain data in that pipeline? (Are the two are complementary to each other?)

Recap

A technique

- Proposed a tailorable sampling pipeline for PVA

A demonstrative use case

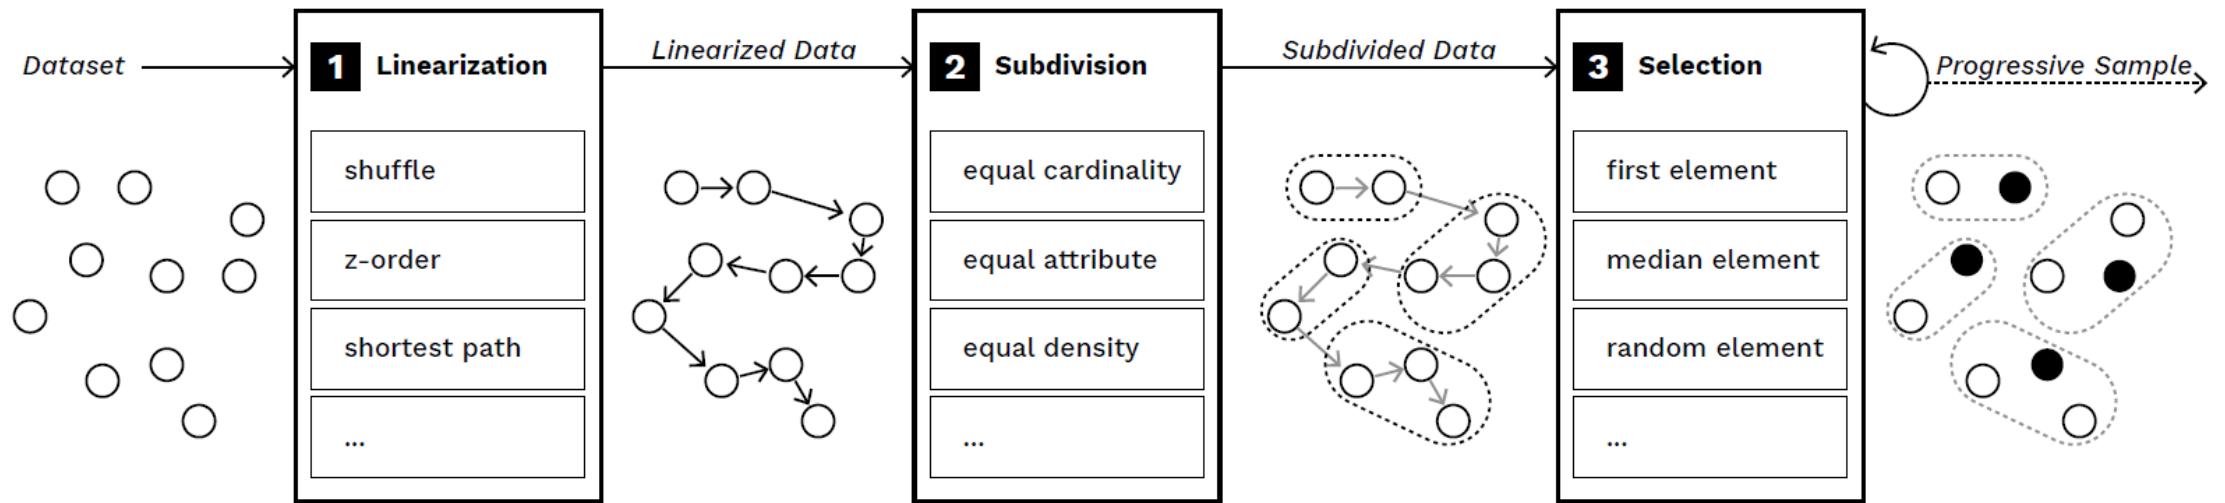
- Showed how to tailor the sampling to three scenarios

A tool

- Introduced ProSample

A Pipeline for Tailored Sampling in Progressive Visual Analytics

Marius Hogräfer, Jakob Burkhardt, Hans-Jörg Schulz



Try ProSample: <https://vis-au.github.io/prosample>